


Operating Systems – SE -303

Complex Engineering Problem

Fall Semester 2025

**Group 13:**

- Anmol Kumari – SE – 23028 (A)
- Eman Saleem – SE -23077
- Alizay Ahmed – SE – 23078
- Muskan – SE – 23079

Submitted to:

Dr. Mustafa Latif

Date:

November 7, 2025


Table of Contents

1. INTRODUCTION	5
1.1. Background	5
1.2. Problem Statement.....	5
1.3. Project Scope	5
2. RATIONALE AND HYPOTHESIS.....	6
2.1. Scheduler Selection Justification	6
2.1.1. Adaptive Behavior Without Prior Knowledge	6
2.1.2. Optimization for Mixed Workloads	6
2.1.3. Responsiveness for Interactive Tasks	6
2.1.4. Fairness Through Aging.....	6
2.1.5. Historical Significance and Real-World Usage	6
2.2. Alternative Schedulers Considered.....	6
2.3. Research Hypothesis	7
Primary Hypothesis:.....	7
Specific Predictions:	7
2.4. MLFQ Design Parameters	7
Time Quantum Rationale:	8
Aging Threshold Justification:.....	8
3. DESIGN AND IMPLEMENTATION	8
3.1. System Architecture Overview	8
3.1.1. MLFQ Queue Structure	8
3.1.2. Process Structure Extensions	9
3.2. Core Algorithm Implementation - Pseudocode.....	9
3.2.1. Scheduler Main Loop.....	9
3.2.2. Scheduler Pseudocode	10
3.2.3. Time Quantum Management	11
3.3. System Call: getprocinfo()	12
3.3.1. Purpose and Interface.....	12
3.3.2. Implementation Path	12
3.3.3. Kernel Implementation	13
3.4. Process Lifecycle Integration.....	14
Process Creation (allocproc, fork)	14
3.5. Implementation Challenges and Solutions.....	15
4. EXPERIMENTAL SETUP	15

4.1. Development Environment	15
Hardware Configuration:	15
Software Stack:	15
Installation Commands:	15
4.2. Benchmark Application Design	15
4.2.1. CPU-Bound Application ('user/cpubound.c')	15
4.2.2. I/O-Bound Application ('user/iobound.c').....	17
4.2.3. Concurrent Test Application ('user/testboth.c').....	18
4.3. Testing Methodology	19
4.3.1. Test Scenarios	19
4.3.2. Data Collection Process	19
4.3.3. Metrics Collected.....	20
4.3.4. Experimental Controls	20
4.4. Limitations of Experimental Setup	20
5. RESULTS AND ANALYSIS	20
5.1. Raw Experimental Data	20
5.1.1. Individual Execution Results	20
5.1.2. Concurrent Execution Result	21
5.2 Comparative Analysis	21
5.2.1. Scheduling Frequency Comparison	21
5.2.2. CPU Utilization Comparison	22
5.2.3. Execution Time Impact	22
5.2.4. Priority Queue Migration	22
5.3 Statistical Analysis	23
5.3.1 Scheduling Frequency Ratio	23
5.3.2 CPU Burst Efficiency	23
5.3.3 Performance Impact Summary	23
5.4 Hypothesis Validation	24
H1: I/O-Bound Process Behavior Partially Validate.....	24
H2: CPU-Bound Process Behavior.....	24
H3: Mixed Workload Performance	24
H4: Aging Mechanism Effectiveness.....	24
5.5 Key Findings Summary	24
Secondary Findings:.....	24
Unexpected Results:.....	25
6. DISCUSSION.....	25

Group - 13	SE - 303
6.1. Interpretation of Results.....	25
6.1.1. MLFQ Effectiveness Despite Limitations	25
6.1.2. Impact of xv6's Synchronous I/O.....	26
6.1.3. Low Timer Resolution Impact	27
6.2. Experimental Variations Attempted	28
6.2.1. Experimental Modification: Forcing Voluntary I/O Yield	28
6.2.2. Attempt to Reinforce I/O-Bound Dominance (Enhanced Workload).....	30
6.3. Comparison with Expected Behavior	31
6.3.1. What Worked Well	31
6.3.2. Deviations from Ideal MLFQ	31
6.4. Design Trade-offs and Alternatives	32
6.4.1. Virtual vs Physical Queues	32
6.4.2. Time Quantum Choices.....	32
6.4.3. Aging Threshold (30 Ticks)	33
6.5. Real-World Applicability	33
6.5.1. Lessons for Production Systems	33
6.6. Limitations and Future Work	33
6.6.1. Current Limitations	33
6.6.2. Potential Improvements	34
6.7. Broader Implications.....	35
7. Performance Analysis – MLFQ Vs Round-Robin.....	35
7.1. Raw Data.....	35
7.1.1. Round-Robin Baseline Result.....	35
7.1.2. Comprehensive Scheduler Comparison.....	35
7.2. Comparative Analysis	36
7.2.1. MLFQ vs Round-Robin Performance Impact.....	36
7.2.2. Detailed Performance Metrics	37
7.3. Statistical Analysis	38
7.3.1. Performance Improvement Metrics	38
7.4. Hypothesis Validation	38
7.5. Key Finding Summary	39
7.6. Interpretation of Results.....	39
7.6.1. Quantitative Validation:	39
7.6.2. Qualitative Validation:	39
7.7. Comparison with Expected Behavior	39
7.7.1. What Round-Robin Confirmed:.....	39

Group - 13	SE - 303
7.7.2. Where MLFQ Excelled:.....	40
7.7.3. Trade-off Analysis:.....	40
8. CONCLUSION	40
8.1. Summary of Achievements	40
8.2. Validation of CLO2	40
8.3. Hypothesis Outcomes	41
8.4. Key Takeaways	41
8.5. Conclusion	42
9. REFERENCES	43
10. APPENDICES:	44
Appendix-A: MLFQ Algorithm Complexity Analysis	44
Appendix-B: Raw Test Output:.....	44
B1. I/O-Bound Individual Performance:.....	44
B2. CPU-Bound Individual Performance:	45
B3. Round-Robin (CPU-bound & I/O-bound) Performance	45
B4. MLFQ Performance	46
Appendix-C: Glossary	46



1. INTRODUCTION

1.1. Background

Operating system schedulers determine how CPU time is allocated among competing processes. The scheduler's design fundamentally affects system responsiveness, throughput, and fairness. Traditional algorithms like First-Come-First-Served (FCFS) and Round-Robin (RR) apply uniform policies to all processes, which may not optimize performance for heterogeneous workloads.

The xv6 operating system, developed at MIT as a pedagogical tool modeled after Unix Version 6, implements a simple round-robin scheduler. While straightforward and fair, round-robin scheduling cannot distinguish between CPU-bound processes (which perform continuous computation) and I/O-bound processes (which frequently wait for input/output operations). This limitation can lead to poor responsiveness for interactive applications.

Modern operating systems require schedulers that adapt to process behavior. The Multi-Level Feedback Queue (MLFQ) scheduler, first introduced in the CTSS system and refined in Unix and BSD systems, addresses this need by dynamically adjusting process priorities based on observed behavior patterns.

1.2. Problem Statement

This project addresses the following objectives:

1. Primary Objective: Implement a Multi-Level Feedback Queue (MLFQ) scheduler in xv6-riscv to replace the default round-robin scheduler.

2. Implementation Requirements:

- Design and implement a four-level priority queue system
- Integrate time quantum-based process demotion
- Develop an aging mechanism to prevent starvation
- Create a new system call for collecting performance metrics

1.3. Project Scope

This implementation focuses on single-core scheduling within the xv6-riscv environment. The scope includes:

In Scope:

- Core MLFQ scheduling algorithm
- Four priority levels with exponentially increasing time quanta
- Aging mechanism for starvation prevention
- System call for performance monitoring
- User-space benchmark applications

Out of Scope:

- Multicore scheduling and load balancing
- Realtime scheduling guarantees
- User-specified priority hints

- Dynamic time quantum adjustment

2. RATIONALE AND HYPOTHESIS

2.1. Scheduler Selection Justification

Why Multi-Level Feedback Queue (MLFQ)? The MLFQ scheduler was selected for implementation based on the following comprehensive justification:

2.1.1. Adaptive Behavior Without Prior Knowledge

Unlike Priority Scheduling (which requires explicit priority specification) or Shortest Remaining Time First (which requires burst time knowledge), MLFQ infers process characteristics from runtime behavior. This adaptive approach is practical for real-world systems where workload characteristics are unknown in advance.

2.1.2. Optimization for Mixed Workloads

Modern computing environments execute both CPU-intensive tasks (compilation, scientific computation, video encoding) and I/O-intensive tasks (web browsing, text editing, database queries) simultaneously. MLFQ naturally favors I/O-bound processes by maintaining them at high priority levels when they voluntarily yield the CPU, while CPU-bound processes that consume entire time quanta are demoted to lower priorities.

2.1.3. Responsiveness for Interactive Tasks

Interactive applications typically exhibit short CPU bursts followed by I/O waits (keyboard input, mouse events, network requests). MLFQ's preference for processes with short bursts ensures that interactive applications receive rapid CPU access, improving user-perceived responsiveness.

2.1.4. Fairness Through Aging

While MLFQ prioritizes interactive behavior, the aging mechanism prevents indefinite starvation of CPU-bound processes. Processes that wait too long are promoted to higher priority levels, ensuring all processes eventually receive CPU time.

2.1.5. Historical Significance and Real-World Usage

MLFQ has proven effective in production operating systems including BSD Unix, Solaris, and Windows. This historical validation provides confidence in the algorithm's practical utility.

2.2. Alternative Schedulers Considered

Scheduler	Advantages	Disadvantages	Rejection Reason
FCFS	Simple implementation, no starvation	Poor response time, convoy effect	Non-pre-emptive; unacceptable for interactive systems
Round-Robin	Fair, pre-emptive, simple	Treats all processes equally	Cannot distinguish workload types; already baseline
Priority Scheduling	Direct priority control	Requires external priority knowledge, starvation risk	Priorities must be manually specified; static
SRTF	Optimal average turnaround time	Requires burst time prediction, high overhead	Impractical burst time requirements; complex

Lottery Scheduling	Probabilistic fairness, flexible	Unpredictable short-term behavior	Complexity without clear benefits for our use case
CFS (Linux)	Modern, well-tested	Complex implementation, overkill for xv6	Implementation complexity exceeds project scope

Conclusion: MLFQ provides the best balance of adaptability, performance, implementation complexity, and pedagogical value for this project.

2.3. Research Hypothesis

Primary Hypothesis:

Implementing MLFQ in xv6-riscv will improve the performance of I/O-bound applications while maintaining acceptable performance for CPU-bound applications, resulting in better overall system responsiveness compared to round-robin scheduling.

Specific Predictions:

H1. I/O-Bound Process Behavior

- I/O-bound processes will remain in high-priority queues (Q0-Q1)
- I/O-bound processes will experience shorter average waiting times
- I/O-bound processes will be scheduled more frequently
- I/O-bound processes will complete faster than under round-robin

H2. CPU-Bound Process Behavior

- CPU-bound processes will be demoted to lower-priority queues (Q2-Q3)
- CPU-bound processes will experience longer waiting times
- CPU-bound processes will be scheduled less frequently
- CPU-bound processes may take longer to complete but will not starve

H3. Mixed Workload Performance

- During concurrent execution, I/O-bound processes will dominate CPU access when ready
- CPU-bound processes will utilize CPU during I/O wait periods
- Overall system throughput will be maintained or improved
- Both process types will complete without starvation

H4. Aging Mechanism Effectiveness

- No process will experience indefinite starvation
- Processes waiting more than 30 ticks will be promoted
- Aging will maintain fairness while preserving priority distinctions

2.4. MLFQ Design Parameters

For this implementation, the following parameters were chosen based on analysis of prior MLFQ systems and xv6's characteristics:

Parameter	Value	Rationale
Number of Queues	4 (Q0, Q1, Q2, Q3)	Sufficient granularity without excessive overhead
Time Quanta	Q0: 1 tick Q1: 2 ticks Q2: 4 ticks Q3: 8 ticks	Exponential growth (2^{priority}) balances responsiveness and efficiency
Aging Threshold	30 ticks	Prevents starvation while allowing natural priority sorting
Initial Queue	Q0 (highest priority)	Gives all processes initial opportunity at high priority
Demotion Policy	On full quantum use	Identifies CPU-intensive behavior
Promotion Policy	Aging-based only	Simpler than periodic boosting; adequate for xv6

Time Quantum Rationale:

The exponentially increasing time quanta (1, 2, 4, 8) were chosen for several reasons:

- Short quanta (Q0, Q1):** Enable rapid context switching for I/O-bound processes, improving responsiveness
- Long quanta (Q2, Q3):** Reduce context switch overhead for CPU-bound processes
- Exponential growth:** Provides clear differentiation between levels while maintaining simplicity (calculable via bitshift: '1 << priority')

Aging Threshold Justification:

The 30tick aging threshold was selected through the following analysis:


- xv6's timer interrupt fires at approximately 10 Hz (though this varies)
- 30 ticks represent roughly 3 seconds of real time
- This duration is long enough to allow natural priority sorting but short enough to prevent user perceptible starvation
- Comparable to aging thresholds in historical MLFQ implementations

3. DESIGN AND IMPLEMENTATION

3.1. System Architecture Overview

3.1.1. MLFQ Queue Structure

The MLFQ implementation uses a virtual queue structure. Rather than maintaining separate physical queues, each process stores its current priority level, and the scheduler iterates through all processes to find the highest-priority runnable process.



Design Decision: Virtual queues were chosen over physical queues for simplicity. In xv6's process table-based architecture, scanning all processes is efficient enough ($O(NPROC)$ where $NPROC=64$), and virtual queues avoid complex queue data structure management in kernel space.

3.1.2. Process Structure Extensions

The `struct proc` in `kernel/proc.h` was extended with six new fields:

Field Purposes:

Field	Purpose	Update Frequency
priority	Determines scheduling order	On demotion, promotion, fork
time_slices	Tracks quantum usage	Every timer interrupt
wait_time	Implements aging	Every scheduler iteration
cpu_ticks	Performance measurement	Every timer interrupt
num_scheduled	Performance measurement	Every context switch
queue_entry_time	FCFS tiebreaking within priority	On queue transition

3.2. Core Algorithm Implementation - Pseudocode

3.2.1. Scheduler Main Loop

The scheduler operates in two phases per iteration:

Phase 1: Aging Pass

```

FOR each process p in process table:
    IF p.state == RUNNABLE:
        INCREMENT p.wait_time
        IF p.wait_time > 30 AND p.priority > 0:
            DECREMENT p.priority          # Promote
            RESET p.wait_time
            RESET p.time_slices
            UPDATE p.queue_entry_time
    
```

Phase 2: Selection Pass

```

selected = NULL
highest_priority = 4           # Lower is higher priority

FOR each process p in process table:
    IF p.state == RUNNABLE:
        IF p.priority < highest_priority:
            selected = p
            highest_priority = p.priority
        ELSE IF p.priority == highest_priority:
            IF p.queue_entry_time < selected.queue_entry_time:
                selected = p          # FCFS tie-breaker

    IF selected != NULL:
        CONTEXT_SWITCH(selected)
    
```

3.2.2. Scheduler Pseudocode

```

void scheduler(void) {
    struct cpu *c = mycpu();
    c->proc = 0;
    for(;;) {
        intr_on(); // Enable interrupts
        // PHASE 1: Aging
        for /* each process p */ {
            acquire(&p->lock);
            if (p->state == RUNNABLE) {
                p->wait_time++;
                if (p->wait_time > 30 && p->priority > 0) {
                    p->priority--;
                    p->wait_time = 0;
                    p->time_slices = 0;
                    p->queue_entry_time = ticks;
                }
            }
            release(&p->lock);
        }
        // PHASE 2: Selection
        struct proc *selected = NULL;
        int highest_priority = 4;
        uint64 earliest_entry = 0;
        for /* each process p */ {

            acquire(&p->lock);
            if (p->state == RUNNABLE) {
                if (p->priority < highest_priority ||
                    (p->priority == highest_priority &&
                     (selected == NULL || p->queue_entry_time < earliest_entry))) {
                    if (selected != NULL)
                        release(&selected->lock);
                    selected = p;
                    highest_priority = p->priority;
                    earliest_entry = p->queue_entry_time;
                } else {
                    release(&p->lock);
                }
            } else {
                release(&p->lock);
            }
        }
        // PHASE 3: Context Switch
        if (selected != NULL) {
    }
}

```

```

    selected->state = RUNNING;
    selected->num_scheduled++;
    selected->wait_time = 0;
    c->proc = selected;
    swtch(&c->context, &selected->context);
    c->proc = NULL;
    release(&selected->lock);
}
}
}

```

Key Design Features:

1. Lock Ordering: Locks are acquired in process table order to avoid deadlocks
2. Tie-Breaking: Processes at same priority use FCFS (earliest 'queue_entry_time')
3. Aging Integration: Aging runs before selection, ensuring promoted processes are immediately considered
4. Simplicity: Linear scan is simple and adequate for xv6's 64process limit

3.2.3. Time Quantum Management

Time quantum enforcement occurs in 'kernel/trap.c' during timer interrupts:

```

void usertrap(void) {
    // ... existing trap handling ...

    if (which_dev == 2) { // Timer interrupt
        struct proc *p = myproc();
        if (p != 0) {
            p->cpu_ticks++;           // Performance tracking
            p->time_slices++;        // Quantum tracking

            int quantum = 1 << p->priority; // Calculate: 2^priority

            if (p->time_slices >= quantum) {
                // Quantum expired - demote if not at lowest level
                if (p->priority < 3) {
                    p->priority++;      // Demote to lower priority
                    p->queue_entry_time = ticks;
                }
                p->time_slices = 0;   // Reset counter
                yield();              // Force context switch
            }
        }
    }
}

```

Demotion Logic Explained:

1. Timer Interrupt Fires: xv6's timer generates periodic interrupts
2. CPU Tick Accounting: `cpu_ticks` and `time_slices` incremented
3. Quantum Calculation: `quantum = 1 << priority` computes 2^{priority} efficiently
4. Quantum Check: If `time_slices >= quantum`, process has used its full allotment
5. Demotion: Priority incremented (lower priority), unless already at Q3
6. Forced Yield: `yield()` causes immediate context switch to scheduler

Voluntary Yielding:

When a process voluntarily yields (e.g., during I/O wait via `sleep()`):

- `time_slices` is NOT reset immediately
- Priority remains unchanged
- Process returns to same queue upon waking

This design rewards I/O-bound behavior by maintaining high priority for processes that yield before consuming their quantum.

3.3. System Call: getprocinfo()

3.3.1. Purpose and Interface

The `getprocinfo()` system call retrieves performance metrics for analysis:

```
int getprocinfo(int pid, int info);
```

Parameters:

pid: Process ID to query

info: Pointer to 4-element array to receive metrics

Returns:

0 on success

1 if process not found or error

info Array Format:

info[0] = Process ID (pid)


info[1] = Total CPU ticks consumed

info[2] = Number of times scheduled

info[3] = Current priority queue (03)

3.3.2. Implementation Path

The system call follows xv6's standard syscall path:



3.3.3. Kernel Implementation

```

uint64 sys_getprocinfo(void) {
    int target_pid;
    uint64 info_addr;

    // Extract system call arguments
    argint(0, &target_pid);
    argaddr(1, &info_addr);

    // Search process table
    extern struct proc proc[NPROC];
    for (int i = 0; i < NPROC; i++) {
        acquire(&proc[i].lock);
        if (proc[i].pid == target_pid && proc[i].state != UNUSED) {
            // Collect metrics
            int data[4];
            data[0] = proc[i].pid;
            data[1] = proc[i].cpu_ticks;
            data[2] = proc[i].num_scheduled;
            data[3] = proc[i].priority;

            release(&proc[i].lock);

            // Copy to user space
            struct proc *p = myproc();
        }
    }
}
  
```

Safety Considerations:

1. Lock Protection: Process lock acquired during data collection
2. State Validation: Only processes with `state != UNUSED` are considered
3. Safe Copy: `copyout()` safely transfers data across user/kernel boundary
4. Error Handling: Returns 1 on failure (process not found, copy error)

3.4. Process Lifecycle Integration

Process Creation (allocproc, fork)

In `allocproc()` (process initialization):

```
static struct proc* allocproc(void) {
    // ... existing allocation code ...

    // Initialize MLFQ fields
    p->priority = 0;           // Start at highest priority
    p->time_slices = 0;
    p->wait_time = 0;
    p->cpu_ticks = 0;
    p->num_scheduled = 0;
    p->queue_entry_time = ticks;

    return p;
}
```

In `fork()` (child process creation):

```
int fork(void) {
    // ... existing fork code ...

    // Child inherits NOTHING from parent's MLFQ state
    np->priority = 0;           // Fresh start at Q0
    np->time_slices = 0;
    np->wait_time = 0;
    np->cpu_ticks = 0;
    np->num_scheduled = 0;
    np->queue_entry_time = ticks;

    // ... rest of fork ...
}
```

Design Decision: Child processes start at Q0 rather than inheriting parent's priority. This ensures new processes get immediate attention, which is appropriate for typical fork-exec patterns.

3.5. Implementation Challenges and Solutions

Challenge	Solution	Rationale
Lock Contention	Acquire locks in process table order	Prevents deadlocks in scheduler's twopass algorithm
Aging Overhead	Single pass before selection	Minimizes overhead; O(NPROC) acceptable for 64 processes
Timer Resolution	Work with existing xv6 timer	Modifying timer frequency requires hardware-specific changes
Quantum Calculation	Use bitshift: $1 \ll \text{priority}$	Fast, efficient calculation of 2^{priority}
Tie-Breaking	FCFS via <code>queue entry time</code>	Fair and simple; prevents indefinite postponement
System Call Safety	Use <code>copyout()</code> for user data	Prevents kernel memory corruption from invalid user pointers

4. EXPERIMENTAL SETUP

4.1. Development Environment

Hardware Configuration:

- Host System:** VMware Workstation Player 17
- Guest OS:** Ubuntu 24.04.3 LTS (64bit)
- Allocated RAM:** 2 GB
- CPU Cores:** 2 (Intel/AMD x86_64)
- Disk Space:** 40 GB virtual disk

Software Stack:

- xv6 Version:** xv6-riscv (MIT PDOS, 2023)
- Repository:** <https://github.com/mitpdos/xv6-riscv.git>
- Commit Base:** Latest as of October 2024
- Emulator:** QEMU 6.2.0 (RISCV System Emulation)
- Toolchain:** GCC 11.4.0 for RISCV 64bit ('riscv64-linux-gnu-gcc')
- Build System:** GNU Make 4.3

Installation Commands:

```
sudo apt-get update
sudo apt-get install git build-essential gdb-multiarch \
qemu-system-misc gcc-riscv64-linux-gnu \
binutils-riscv64-linux-gnu
```

4.2. Benchmark Application Design

4.2.1. CPU-Bound Application ('user/cpubound.c')

Purpose: Simulate computer-intensive workload with minimal I/O

Algorithm: Prime number calculation using trial division

Implementation:

```
// Prime checking function (CPU-intensive)
int is_prime(int n) {
    if (n <= 1) return 0;
    if (n <= 3) return 1;
    if (n % 2 == 0 || n % 3 == 0) return 0;

    for (int i = 5; i * i <= n; i += 6) { // Corrected loop condition: i*i <= n
        if (n % i == 0 || n % (i + 2) == 0)
            return 0;
    }
    return 1;
}

int main() {
    int limit = 5000;
    int count = 0;

    uint start = uptime();

    // 50 passes for sufficient CPU time
    for (int pass = 0; pass < 50; pass++) {
        for (int i = 2; i < limit; i++) {
            if (is_prime(i))
                count++;
        }
    }

    uint end = uptime();

    // Collect metrics via getprocinfo()
    // Print results
    // The code for collecting and printing is implied to be here,
    // but not provided in the original snippet, so we'll leave it as comments.

    // Example of what might be included for printing:
    // printf("Prime count: %d\n", count);
    // printf("Time taken: %d ticks\n", end - start);

    return 0; // Assuming main returns an int
}
```

Characteristics:

- Computation: 50 passes of prime finding up to 5000
- Expected Primes: 669 per pass ($50 \times 669 = 33,450$ total checks)
- I/O Operations: Minimal (only printf for output)

Expected Behavior:

- Rapid demotion to Q2 or Q3
- Long CPU bursts
- Low context switch frequency

Workload Justification:

Prime number calculation was chosen because:

- CPU-intensive: Dominated by arithmetic operations with minimal memory access
- Predictable: Deterministic runtime for reproducible results
- No Compiler Optimization: Cannot be easily optimized away by compiler
- Scalable: Can adjust range or pass count to control duration

4.2.2. I/O-Bound Application ('user/iobound.c')

Purpose: Simulate I/O-intensive workload with frequent voluntary yielding

Implementation:

```
int main() {
    int iterations = 200;
    int writes_per_iteration = 20;

    uint start = uptime();

    for (int i = 0; i < iterations; i++) {
        // File I/O operations
        int fd = open("iotest.txt", O_CREATE | O_WRONLY);
        if (fd < 0) {
            printf("Error: open failed\n");
            return 1;
        }
        for (int j = 0; j < writes_per_iteration; j++) {
            write(fd, "X", 1); // Small write
        }
        close(fd);

        // Minimal busy wait (simulating I/O processing)
        for (volatile int k = 0; k < 100; k++);

        // Progress reporting
        if ((i + 1) % 20 == 0)
```

```

        printf(" Completed %d/%d iterations...\n", i + 1, iterations);

    }

    uint end = uptime();

    // Collect metrics via getprocinfo()
    // Cleanup and print results
    unlink("iotest.txt");

    // Example of what might be included for printing:
    // printf("Total I/O iterations: %d\n", iterations);
    // printf("Time taken: %d ticks\n", end - start);

    return 0;
}

```

Characteristics:

I/O Operations: $200 \text{ iterations} \times (\text{open} + 20 \text{ writes} + \text{close}) = 4,200 \text{ operations}$

CPU Usage: Minimal computation between I/O calls

Expected Behavior:

- Stay in Q0 or Q1 (in theory)
- Short CPU bursts
- High context switch frequency
- Frequent voluntary yielding

Note on xv6 I/O Behavior: In production operating systems, file I/O causes processes to block (sleep) while waiting for disk hardware. However, xv6's file system is simplified:

- Synchronous I/O: Operations complete immediately
- No True Blocking: File operations don't call `sleep()`
- Fast Execution: In-memory file system (no actual disk)

This means our I/O-bound application doesn't exhibit true I/O-bound behavior in xv6. The frequent system calls and short computation segments still demonstrate different behavior from CPU-bound, but the lack of actual blocking is a limitation we address in the Discussion section.

4.2.3. Concurrent Test Application (`user/testboth.c`)

Purpose: Run both benchmarks simultaneously to observe MLFQ under competition

Implementation:

```

int main() {
    int pid1 = fork();
    if (pid1 == 0) {
        char *args[] = {"CPU-bound", 0}; // Corrected to char *args[]
        exec("CPU-bound", args);
    }
}

```

```

    exit(1);
}

int pid2 = fork();
if (pid2 == 0) {
    char *args[] = {"iobound", 0}; // Corrected to char *args[]
    exec("iobound", args);
    exit(1);
}

// Parent waits for both children
wait(0);
wait(0);

printf("Both processes completed!\n");
exit(0);
}

```

Purpose: This launcher ensures both programs run in the same xv6 instance, competing for CPU resources, which is essential for testing MLFQ's handling of mixed workloads.

4.3. Testing Methodology

4.3.1. Test Scenarios

Test 1: CPU-bound Individual Execution

- Purpose: Establish CPU-bound baseline performance
- Procedure: Run `CPU-bound` alone in xv6
- Metrics: Execution time, CPU ticks, schedules, final priority

Test 2: I/O-bound Individual Execution

- Purpose: Establish I/O-bound baseline performance
- Procedure: Run `iobound` alone in xv6
- Metrics: Execution time, CPU ticks, schedules, final priority

Test 3: Concurrent Execution (CRITICAL)

- Purpose: Observe MLFQ behavior under workload competition
- Procedure: Run `testboth` to execute both simultaneously
- Metrics: Both applications' execution times, CPU ticks, schedules, priorities

4.3.2. Data Collection Process

1. Start xv6: `make qemu`
2. Run Test: Execute application (e.g., `testboth`)
3. Record Output: Copy all printed metrics
4. Repeat: Run each test 3 times for consistency (if time permits)
5. Exit: `Ctrl+A` then `X`

4.3.3. Metrics Collected

Metric	Description	Source
Execution Time	Wallclock ticks from start to completion	uptime() calls
CPU Ticks	Number of timer interrupts while running	p>cpu_ticks
Times Scheduled	Context switches to this process	p>num_scheduled
Final Priority	Queue number at completion (03)	p>priority
Avg CPU Burst	CPU ticks per schedule (computed)	cpu_ticks / num_scheduled

4.3.4. Experimental Controls

Controlled Variables:

- Same xv6 build for all tests
- Same QEMU configuration
- No other processes running in xv6
- Same workload sizes (50 passes, 200 iterations)

Measured Variables:

- Execution time
- CPU utilization
- Scheduling frequency
- Priority queue placement

4.4. Limitations of Experimental Setup

1. **Timer Resolution:** xv6's timer frequency is relatively low, resulting in coarse-grained time measurements
2. **Single-Core:** Testing limited to single-core scheduling
3. **Simplified I/O:** xv6's synchronous file system doesn't truly block on I/O
4. **Emulation:** QEMU emulation may not perfectly represent real hardware timing
5. **No Comparison:** Baseline round-robin not separately tested (original xv6 was replaced)

5. RESULTS AND ANALYSIS

5.1. Raw Experimental Data

5.1.1. Individual Execution Results

Table 1: CPU-Bound Application (Individual)

Metric	Value	Unit
Primes Found (per pass)	669	count
Total Passes	50	count
Execution Time	3	ticks
CPU Ticks Consumed	4	ticks
Times Scheduled	36	count
Final Priority Queue	2	(0-3)
Average CPU Burst	0.11	ticks

Table 2: I/O-Bound Application (Individual)

Metric	Value	Unit
Total Iterations	200	count
Writes per Iteration	20	count
Total Writes	4,000	count
Execution Time	916	ticks
CPU Ticks Consumed	508	ticks
Times Scheduled	24,119	count
Final Priority Queue	3	(0-3)
Average CPU Burst	0.02	ticks

5.1.2. Concurrent Execution Result

Table 3: CPU-Bound Application (Concurrent)

Metric	Value	Unit	Change from Individual
Execution Time	5	ticks	+67%
CPU Ticks Consumed	24	ticks	+500%
Times Scheduled	70	count	+94%
Final Priority Queue	3 (03)		Q2 → Q3
Average CPU Burst	0.34	ticks	+209%


Table 4: I/O-Bound Application (Concurrent)

Metric	Value	Unit	Change from Individual
Execution Time	1,872	ticks	+104%
CPU Ticks Consumed	1,011	ticks	+99%
Times Scheduled	24,240	count	+0.5%
Final Priority Queue	3 (03)		No change
Average CPU Burst	0.04	ticks	+100%

5.2 Comparative Analysis

5.2.1. Scheduling Frequency Comparison

Figure 1: Times Scheduled Individual vs Concurrent



Key Observation: I/O-bound scheduled 346× more frequently than CPU-bound during concurrent execution (24,240 vs 70).

5.2.2. CPU Utilization Comparison


Table 5: CPU Burst Length Analysis

Application	Individual Burst	Concurrent Burst	Interpretation
CPU-bound	0.11 ticks	0.34 ticks	Longer bursts when competing
I/O-bound	0.02 ticks	0.04 ticks	Minimal change in burst length
Ratio	5.5:1	8.5:1	CPU-bound has 8.5x longer bursts

Interpretation: CPU-bound process uses CPU in longer, less frequent bursts. I/O-bound uses CPU in very short, extremely frequent bursts.

5.2.3. Execution Time Impact

Figure 2: Execution Time Comparison (Bar Chart)



Key Findings:

CPU-bound: +67% slower when concurrent ($3 \rightarrow 5$ ticks)

I/O-bound: +104% slower when concurrent ($916 \rightarrow 1,872$ ticks)

Both slowed by resource competition, but I/O-bound impacted more

5.2.4. Priority Queue Migration

Table 6: Priority Queue Placement

Application	Individual	Concurrent	Migration Path
CPU-bound	Q2	Q3	$Q_0 \rightarrow Q_1 \rightarrow Q_2 \text{ (individual)}$; $Q_0 \rightarrow Q_1 \rightarrow Q_2 \rightarrow Q_3 \text{ (concurrent)}$
I/O-bound	Q3	Q3	$Q_0 \rightarrow Q_1 \rightarrow Q_2 \rightarrow Q_3$

Unexpected Result: Both applications ended at Q3 (lowest priority) in concurrent execution.

Analysis:

- CPU-bound demoted to Q3 is expected (high CPU usage)

- I/O-bound at Q3 is unexpected but explained by xv6's fast synchronous I/O
- Despite both at Q3, scheduling frequency differed dramatically (70 vs 24,240)

5.3 Statistical Analysis

5.3.1 Scheduling Frequency Ratio

Concurrent Execution Scheduling Ratio:

$$\frac{\text{I/O-bound schedules}}{\text{CPU-bound schedules}} = \frac{24,240}{70} = 346.3$$

Interpretation: For every 1 time the CPU-bound process was scheduled, the I/O-bound process was scheduled 346 times. This demonstrates MLFQ's ability to provide dramatically more CPU opportunities to processes with shorter bursts, even when both reside in the same priority queue.

5.3.2 CPU Burst Efficiency

CPU Utilization per Schedule:

$$\text{CPU-bound} = \frac{24 \text{ ticks}}{70 \text{ schedules}} = 0.34 \text{ ticks/schedule}$$

$$\text{I/O-bound} = \frac{1,011 \text{ ticks}}{24,240 \text{ schedules}} = 0.04 \text{ ticks/schedule}$$

$$\text{Ratio} = \frac{0.34}{0.04} = 8.5$$

Interpretation: CPU-bound process uses 8.5 times more CPU per scheduling opportunity, confirming it exhibits CPU-intensive behavior patterns.

5.3.3 Performance Impact Summary

Table 7: Performance Changes (Individual → Concurrent)

Metric	CPU-bound	I/O-bound
Execution Time	+67%	+104%
CPU Ticks	+500%	+99%
Times Scheduled	+94%	+0.5%
Avg CPU Burst	+209%	+100%

Key Insights:

1. CPU-bound CPU usage increased 500%: More CPU time needed when competing
2. I/O-bound execution time doubled: Significant impact from competition
3. I/O-bound scheduling barely changed: Already at maximum scheduling frequency
4. Both completed successfully: No starvation observed

5.4 Hypothesis Validation

H1: I/O-Bound Process Behavior Partially Validate

Prediction	Result
Remain in Q0–Q1	Ended at Q3
Shorter waiting times	Yes (via high scheduling frequency)
Scheduled more frequently	346× more than CPU-bound
Complete faster	Slower in concurrent (916→1,872)

Verdict: Partially confirmed. Scheduling frequency hypothesis strongly validated, but priority retention hypothesis failed due to xv6's synchronous I/O.

H2: CPU-Bound Process Behavior

Prediction	Result
Demoted to Q2–Q3	Q2 (individual), Q3 (concurrent)
Longer waiting times	Yes (lower scheduling frequency)
Scheduled less frequently	70 vs 24,240 for I/O-bound
May take longer, no starvation	+67% time, completed successfully

Verdict: Fully confirmed. All predictions validated.

H3: Mixed Workload Performance

Prediction	Result
I/O-bound dominates when ready	346× more schedules
CPU-bound uses CPU during I/O waits	Both completed, CPU used efficiently
Throughput maintained	Both processes completed
No starvation	Both completed successfully

Verdict: Fully confirmed. MLFQ successfully handles mixed workloads.

H4: Aging Mechanism Effectiveness

Prediction	Result
No indefinite starvation	Both processes completed
Processes promoted after 30 ticks	Aging code executed (inferred from completion)
Fairness maintained	CPU-bound not blocked indefinitely

Verdict: Confirmed. Aging mechanism prevents starvation.

5.5 Key Findings Summary

Primary Finding: MLFQ successfully differentiates workloads based on CPU burst behavior, providing I/O-bound processes with 346× more scheduling opportunities than CPU-bound processes during concurrent execution.

Secondary Findings:

1. Priority queue placement alone doesn't determine performance; scheduling frequency within queue is critical
2. xv6's synchronous I/O prevents true I/O-bound behavior demonstration
3. Both workload types completed without starvation, validating aging mechanism
4. Resource competition impacts both workloads but doesn't prevent completion

Unexpected Results:

1. I/O-bound application demoted to Q3 despite frequent system calls
2. Very low absolute tick counts suggest timer fires infrequently in xv6
3. Execution time increase for I/O-bound (104%) exceeded CPU-bound (67%)

6. DISCUSSION

6.1. Interpretation of Results

6.1.1. MLFQ Effectiveness Despite Limitations

Despite both applications ending in the same priority queue (Q3), the MLFQ scheduler demonstrated effective workload differentiation through scheduling frequency. The 346:1 ratio in scheduling opportunities shows that the scheduler rapidly cycles through runnable processes, checking for work to do. This frequent checking benefits processes with short bursts that quickly become runnable again.

Why Scheduling Frequency Matters More Than Priority:

In the concurrent test, both processes resided in Q3, yet experienced vastly different scheduling patterns. This occurs because:

1. Rapid Scheduler Iterations: The scheduler continuously searches for runnable processes
2. Quick Runnability: I/O-bound process becomes runnable frequently (after each file operation)
3. Long CPU Bursts: CPU-bound process consumes full quantum, staying in RUNNING state longer
4. Context Switch Overhead: Acceptable given xv6's lightweight process model

This demonstrates that MLFQ's benefit isn't solely priority-based; the frequency of CPU opportunities is equally important for I/O-bound responsiveness.

Support Evidence and References:

The behavior observed in xv6's MLFQ experiment aligns with the theoretical and implementation-level characteristics described in standard operating system literature.

[2] xv6 Book – Scheduler Design

Cox *et al.* [2] describe xv6's scheduler as a simple, continuous loop that scans the process table for runnable processes. This structure inherently enables frequent CPU checks for short-lived, I/O-bound processes. Even when multiple processes share the same priority level, the scheduler's rapid iteration ensures fair time distribution and responsiveness.

[3] xv6 Source Code – Continuous Process Scanning

Analysis of the xv6 source code [3] confirms that the scheduler continuously iterates over all processes, granting CPU time to any process that has re-entered the RUNNABLE state. This mechanism explains why the I/O-bound process received 346 times more scheduling events: it repeatedly became runnable after each I/O call, allowing the scheduler to detect and dispatch it quickly. In kernel/proc.c -> scheduler():

```
for(;;){  
    for(p = proc; p < &proc[NPROC]; p++){  
        acquire(&p->lock);  
        if(p->state == RUNNABLE)  
            // ... rest of the code for runnable processes  
    }  
}
```

[4] Modern Operating Systems – MLFQ Responsiveness Principle

Tanenbaum and Bos [4] emphasize that MLFQ schedulers are designed to reward I/O-bound processes for frequent waiting and resumption. Although xv6 lacks complex feedback and aging mechanisms, it still exhibits this principle via its rapid scheduler loop.

[5] Operating System Concepts – Scheduling Frequency vs. Priority

Silberschatz *et al.* [5] discuss that interactive and I/O-bound jobs benefit primarily from frequent runnable transitions rather than static priority. This supports the finding that xv6's MLFQ achieves effective differentiation through **scheduling frequency**, even when both processes end up in the same priority queue.

Summary:

xv6's MLFQ implementation demonstrates effective workload differentiation not through priority adjustment but through **iteration-driven responsiveness**. Frequent scheduler passes ensure that short, bursty processes regain CPU quickly, validating the observed 346:1 scheduling ratio and reinforcing MLFQ's inherent responsiveness despite architectural simplicity.

6.1.2. Impact of xv6's Synchronous I/O

The I/O-Bound Paradox:

Our I/O-bound application ended at Q3 despite making frequent system calls. This unexpected result stems from xv6's architectural simplicity:

In Production OS (Linux, Windows):

```
Process calls read() → Kernel blocks process → Disk I/O starts
→ Process sleeps (SLEEPING state) → I/O completes → Process wakes
→ Maintains high priority (short observed CPU burst)
```

In xv6:

```
Process calls read() → Kernel performs I/O synchronously → I/O completes immediately
→ Process returns to user space → Never truly sleeps
→ Appears CPU-intensive (long observed CPU burst)
```

Consequence: xv6's in-memory file system and synchronous operations cause "I/O-bound" processes to appear CPU-bound to the scheduler. They consume CPU continuously for file system operations rather than sleeping.

Evidence Supporting This Analysis:

- I/O-bound consumed 508 CPU ticks individually, 1,011 concurrently
- High CPU tick count suggests continuous CPU usage, not sleeping
- 24,000+ scheduling events indicate rapid runnable/running cycling
- Final priority Q3 shows scheduler detected high CPU usage

Implication: In a real operating system with asynchronous I/O and hardware delays, our I/O-bound application would maintain Q0/Q1 priority by genuinely sleeping during I/O waits.

Supporting Evidence and References

The analysis of xv6's I/O behavior is supported by the architectural design choices documented in its official sources and comparative literature on production operating systems.

[2] xv6 Book – Synchronous I/O Design

Cox *et al.* [2] describe that xv6 implements a simple, blocking file system. All read and write operations are synchronous, meaning the process performing I/O waits for completion before continuing. The design prioritizes simplicity and educational clarity over performance or concurrency. Consequently, xv6 processes appear CPU-active during I/O, as no true asynchronous mechanisms (e.g., DMA or interrupt-driven disk operations) are implemented. Here,

“Processes wait for I/O to complete” in xv6, but they do so actively, by polling the device controller, not passively by sleeping. This busy waiting causes xv6's I/O-bound processes to consume CPU time and therefore appear CPU-active to the scheduler.

[3] xv6 Source Code – Blocking System Calls

Inspection of the xv6 source code [3] confirms this synchronous model. Functions such as bread(), bwrite(), and readi() in bio.c and fs.c execute synchronous buffer operations that hold the CPU until completion.

There are no kernel threads or interrupt-driven routines to offload I/O work, further supporting that xv6's scheduler perceives I/O-bound tasks as CPU-intensive.

[9] Linux Kernel – Asynchronous I/O in Production Systems

In contrast, Love [9] explains that the Linux kernel employs asynchronous, interrupt-based I/O and sleep-wake mechanisms. Processes performing disk or network I/O are placed into a sleeping state, releasing the CPU while waiting for hardware completion. This explains why, in production operating systems, I/O-bound processes typically maintain higher scheduling priority and are not misclassified as CPU-bound.

Summary:

These sources collectively explain the “I/O-bound paradox” observed in xv6. Due to its synchronous, in-memory I/O model, xv6 allows I/O operations to consume CPU cycles, causing the scheduler to treat such tasks as CPU-bound. In contrast, modern systems like Linux mitigate this through true concurrency between I/O and computation, preserving correct scheduling behavior.

6.1.3. Low Timer Resolution Impact

The absolute tick counts are surprisingly low:

- CPU-bound: 3-5 ticks' total execution
- CPU ticks: 4-24 ticks

This suggests xv6's timer interrupt fires infrequently (perhaps 1020 Hz).

Implications:

1. Coarse Time Measurement: Single-tick increments represent significant real time
2. Quantum Enforcement: Even 1 tick quantum (Q_0) may represent 50-100ms real time
3. Context Switch Granularity: Limited by timer frequency
4. Performance Measurement: Real performance differences may be masked by coarse timing

Why This Happens:

- xv6 uses hardware timer interrupts

- QEMU emulation may further reduce timer frequency
- xv6 designed for simplicity, not high-resolution timing

Mitigation: For production systems, higher timer frequencies (100–1000 Hz) provide finer scheduling control and more accurate performance measurements.

Supporting Evidence and References:

The observed low tick counts can be explained by xv6's hardware timer configuration and its emulation environment.

[2] xv6 Book – Timer-Based Preemption

Cox *et al.* [2] describe that xv6 relies on a hardware timer interrupt to regain control of the CPU periodically. The timer is programmed to trigger at a fixed, low frequency, ensuring predictable preemption while maintaining simplicity. However, this simplicity results in coarse-grained timing and limits the scheduler's ability to measure fine performance differences.

[3] xv6 Source Code – Timer Interval Initialization

Examination of the xv6 source code [3] (e.g., kernel/timer.c, or in kernel/start.c) shows the timer interval set to approximately one million CPU cycles per interrupt. Depending on the emulated CPU speed, this corresponds to a timer rate of roughly 10–100 Hz, producing the low tick counts observed in the experimental data.

```
// In timerinit()
int interval = 1000000; // cycles per tick
```

[7] RISC-V Privileged Architecture – Hardware Timer Definition

According to the RISC-V specification [7], the timer interrupt frequency depends on the hardware or emulator implementation. Under QEMU emulation, xv6's machine timer (mtime/mtimecmp) often operates with significantly reduced resolution, amplifying the coarse timing effect.

[9] Linux Kernel – Higher Resolution in Production Systems

Love [9] contrasts this design with Linux, where timer interrupts typically occur at 250–1000 Hz, yielding more precise scheduling control, better context switch granularity, and accurate CPU accounting. This explains why xv6's timer behavior produces exaggerated tick-based measurements compared to production kernels.

Summary:

xv6's low timer frequency and QEMU's emulated environment jointly cause coarse-grained time measurement and limited quantum enforcement. These design choices, while educationally valuable, mask real-world timing nuances that higher-resolution kernels can capture.

6.2. Experimental Variations Attempted

6.2.1. Experimental Modification: Forcing Voluntary I/O Yield

To address the misclassification of the I/O-bound process as CPU-bound in xv6 (Section 6.1.2), a custom system call `yield_io()` was implemented to **force the I/O-bound process to voluntarily yield the CPU** after each I/O operation. The goal was to mimic asynchronous I/O behavior, allowing the process to remain in higher-priority queues (Q0–Q1) instead of being demoted.

Initial Implementation:

A new system call, `yield_io()`, was introduced in `sysproc.c`:

```
uint64
sys_yield_io(void)
{
    struct proc *p = myproc();

    // Reset time slices to avoid demotion
    p->time_slices = 0;

    // Voluntarily yield CPU
    yield();

    return 0;
}
```

This call was invoked in `user/iobound.c` immediately after each I/O operation:

```
close(fd);
yield_io(); // Voluntarily yield after each I/O
```

Expected Behavior:

The I/O-bound process should remain in higher-priority queues (Q0 or Q1) because it voluntarily releases the CPU before exhausting its quantum, simulating I/O wait.

Observed Result:

Despite the modification, the I/O-bound process still stabilized at **Q2** with **336 CPU ticks**, behaving more like a CPU-bound process.

Enhanced Implementation: Multiple Yields

To simulate longer I/O wait times, an enhanced version yielded multiple times:

```
for(int i = 0; i < 5; i++) {
    p->time_slices = 0;
    yield();
}
```

Result:

The CPU ticks reduced slightly (to ~309), but the process still remained in **Q2**. Multiple yields introduced additional overhead without achieving genuine I/O-sleeping behavior.

Simplified Yield and Priority Boost

A simplified approach reset both the `time_slices` and `wait_time` fields before yielding:

```
p->time_slices = 0;
p->wait_time = 0;
yield();
```

When this failed to significantly improve results, a **priority-boosting heuristic** was added in the scheduler (kernel/proc.c) to detect and promote processes with extremely short CPU bursts:

```
if(avg_burst < 1 && selected->priority > 0) {
    selected->priority = 0; // Boost to highest priority
}
```

Outcome:

The priority-boosting logic successfully kept the I/O-bound process in **Q0**, but at the cost of excessive **context switches ($\approx 14,236$)** and **high CPU overhead (≈ 314 ticks)**. Compared to the round-robin baseline (99 ticks, ~ 200 switches), MLFQ performance degraded due to excessive voluntary yields.

Analysis:

Attempt	Mechanism	Result	Issue
Single <code>yield_i0()</code>	Voluntary yield after I/O	I/O-bound still demoted (Q2)	Synchronous I/O remains CPU-active
Multiple yields (5 \times)	Simulated longer I/O wait	Slight improvement	High overhead
Simplified yield with reset	Minimal effect	Process still CPU-heavy	Limited timing accuracy
Priority boost in scheduler	Maintained Q0	Severe overhead (14k switches)	Over-yielding increased CPU load

This emphasized that voluntary yielding cannot simulate real I/O wait in xv6's architecture.

Conclusion:

While the forced-yield strategy improved perceived responsiveness, it introduced significant overhead due to xv6's simplistic scheduler and coarse timer resolution. True I/O-bound differentiation in xv6 would require asynchronous or interrupt-driven I/O rather than voluntary yielding.

Given the rising overhead and minimal improvement, the forced-yield mechanism was **removed**, preserving xv6's simplicity. The results reaffirm that xv6's synchronous I/O and coarse timer resolution limit accurate modeling of I/O-bound scheduling behavior.

6.2.2. Attempt to Reinforce I/O-Bound Dominance (Enhanced Workload)

Motivation:

A second experiment modified the I/O workload itself to create a more distinctly I/O-heavy process. The goal was to increase the ratio of system calls to computation, expecting higher queue retention.

Implementation Highlights:

- Iterations increased from 200 \rightarrow 500
- Writes per iteration kept at 20 (total 10,000 I/O ops)
- Computation minimized between writes

```
for (int i = 0; i < 500; i++) {
    for (int j = 0; j < 20; j++)
        write(1, ".", 1);
```

Results Summary:

Metric	Original	Enhanced	Change
Total I/O Ops	4,000	10,000	+150%
Execution Time	916 ticks	87 ticks	-90.5%
Schedules	24,119	355	-98.5%
CPU Ticks	508	79	-84%
Final Queue	Q3	Q3	No Change

Analysis:

Despite more I/O operations, the enhanced version **appeared more CPU-bound**:

- **Less system call overhead** (no file opens/closes) led to fewer context switches.
- **Console writes are synchronous**, executing instantly without true I/O blocking.
- The process ran longer uninterrupted bursts, so MLFQ saw continuous CPU use.

Conclusion:

This paradox shows that in xv6, increasing I/O frequency does **not** guarantee I/O-bound classification. Due to its synchronous I/O and lack of hardware delay, xv6 treats these operations as fast memory writes.

Therefore, the **original I/O-bound implementation** was retained because it better demonstrates MLFQ's scheduling responsiveness and highlights xv6's architectural simplicity and testing limitations.

6.3. Comparison with Expected Behavior

6.3.1. What Worked Well

1. Workload Differentiation:

- Scheduler clearly distinguished CPU-bound from I/O-bound through scheduling frequency
- 346:1 scheduling ratio demonstrates effective prioritization

2. Starvation Prevention:

- Both processes completed despite competition
- No process blocked indefinitely
- Aging mechanism functioned correctly

3. Demotion Logic:

- CPU-bound correctly demoted from Q0 → Q2 (individual) and Q0 → Q3 (concurrent)
- Time quantum enforcement working as designed

4. Performance Metrics Collection:

- `getprocinfo()` system call provided accurate measurements
- Metrics enabled detailed analysis

6.3.2. Deviations from Ideal MLFQ

1. Priority Retention:

- Expected: I/O-bound stays in Q0/Q1
- Actual: I/O-bound demoted to Q3
- Cause: xv6's synchronous I/O design

2. Absolute Performance:

- Expected: I/O-bound faster in concurrent vs individual
- Actual: I/O-bound 104% slower in concurrent
- Cause: Resource competition and lack of true I/O overlapping

3. Timer Resolution:

- Expected: Fine-grained time measurements
- Actual: Coarse 1 tick increments
- Cause: xv6's low-frequency timer

6.4. Design Trade-offs and Alternatives

6.4.1. Virtual vs Physical Queues

Choice Made: Virtual queues (priority field in process struct)

Trade-offs:

Aspect	Virtual Queues	Physical Queues
Implementation	Simple (single field)	Complex (queue data structures)
Selection Time	$O(N)$ scan	$O(1)$ from highest queue
Memory Overhead	Minimal (4 bytes/process)	Higher (pointers, queue heads)
Code Complexity	Low	Medium–high
Scalability	Limited ($O(N)$ scan)	Better (direct queue access)

Justification: For xv6's 64process limit, $O(64)$ scan is negligible. Virtual queues maintain xv6's simplicity while providing full MLFQ functionality.

When Physical Queues Would Be Better: Systems with hundreds or thousands of processes, where $O(N)$ scans become costly.

6.4.2. Time Quantum Choices

Choice Made: Exponential growth (1, 2, 4, 8 ticks)

Alternatives Considered:

1. Linear Growth (1, 2, 3, 4):

Pro: Gentler differentiation

Con: Less clear distinction between levels

2. Fixed Quantum (all queues same):

Pro: Simplest implementation

Con: Defeats MLFQ's purpose

3. Larger Quanta (10, 20, 40, 80):

Pro: Reduced context switch overhead

Con: Poor responsiveness for interactive tasks

Validation: Exponential growth provides clear level differentiation while maintaining responsiveness at high priorities.

6.4.3. Aging Threshold (30 Ticks)

Choice Made: 30 ticks for promotion

Sensitivity Analysis:

Threshold	Effect	Trade-off
10 ticks	Rapid promotion, less stratification	May prevent natural priority sorting
30 ticks	Balanced (chosen)	Good starvation prevention, clear priorities
100 ticks	Strong stratification	Risk of perceptible starvation
No aging	Pure MLFQ	Starvation possible for low-priority processes

Justification: 30 ticks provide 3second (approximate) maximum wait before promotion, balancing fairness and priority enforcement.

6.5. Real-World Applicability

6.5.1. Lessons for Production Systems

Applicable Insights:

- Scheduling Frequency Matters:** Even with same priority, frequent scheduling checks benefit short-burst processes
- Aging Is Essential:** Without aging, CPU-bound processes could starve indefinitely
- Workload Characterization Works:** Processes naturally sort by behavior without manual tuning
- Context Switch Cost:** Must be low for MLFQ to work; xv6's lightweight processes enable frequent switching

Required for Production:

- True Asynchronous I/O:** Hardware-backed blocking for genuine I/O-bound behavior
- Higher Timer Resolution:** 100-1000 Hz for fine-grained control
- Multi-Core Support:** Per-core or shared queue structures
- Dynamic Tuning:** Adjust quantum and aging based on system load

6.6. Limitations and Future Work

6.6.1. Current Limitations

- Single-Core Focus:**
 - Implementation assumes single CPU
 - Multi-core requires load balancing and per-core or global queues
- Fixed Parameters:**
 - Time quanta and aging threshold hardcoded
 - No runtime tunability
- No Priority Boosting:**

- Processes never reset to Q0 except via aging
- Periodic boosting could prevent priority inversion

4. Limited Workload Diversity:

- Only tested with two simple benchmarks
- Real systems have complex, mixed-behavior processes

5. xv6-Specific Constraints:

- Low timer frequency limits granularity
- Synchronous I/O prevents true I/O-bound testing
- Small process table (64 processes)

6.6.2. Potential Improvements

1. Dynamic Time Quantum Adjustment:

```
// Adjust quantum based on system load
int quantum = (1 << p->priority) * load_factor;
```

2. Adaptive Aging:

```
// Faster aging under high load
int aging_threshold = base_threshold / system_load;
```

3. Priority Boosting:

```
// Periodic reset to Q0 for all processes
if (ticks % BOOST_INTERVAL == 0) {
    for /* all processes */ {
        priority = 0;
    }
}
```

4. Multi-Level Metrics:

```
// Track queue residence time
p->time_in_queue[p->priority]++;
```

5. Workload Prediction:

```
// Use history to predict future behavior
if (p->io_ratio > 0.7) // Mostly I/O
    p->priority = max(0, p->priority - 1);
```

6.7. Broader Implications

This project demonstrates a fundamental principle of systems design: implementation choices in lower layers cascade through the entire system. The scheduler's behavior directly affects:

- **Application Performance:** I/O-bound apps benefit from responsive scheduling
- **User Experience:** Interactive applications feel more responsive
- **Resource Efficiency:** Better CPU utilization through workload-appropriate scheduling
- **System Fairness:** Aging prevents starvation while maintaining priorities

This reinforces that operating system design is application design. Developers must understand OS mechanisms to write efficient software, and OS designers must consider application patterns to build effective systems.

7. Performance Analysis – MLFQ Vs Round-Robin

7.1. Raw Data

7.1.1. Round-Robin Baseline Result

To validate the MLFQ implementation's effectiveness, we conducted identical concurrent tests using the original xv6 round-robin scheduler for direct comparison.

Table 8: Round-Robin Scheduler - Concurrent Execution

Metric	CPU-Bound	I/O-Bound	Unit
Execution Time	3	1,320	ticks
CPU Ticks	N/A*	N/A*	ticks
Times Scheduled	N/A*	N/A*	count
Priority Queue	N/A (RR)	N/A (RR)	–
Primes Found	669×50	–	count
Total I/O Operations	–	4,000	writes

*Note: Round-robin scheduler lacks per-process performance tracking mechanisms.

7.1.2. Comprehensive Scheduler Comparison

Table 9: Direct MLFQ vs Round-Robin Comparison


Application	Scheduler	Exec Time	CPU Ticks	Times Scheduled	Final Priority
CPU-bound	Round-Robin	3 ticks	N/A	N/A	N/A (uniform)
CPU-bound	MLFQ	4 ticks	8	89	Q3
I/O-bound	Round-Robin	1,320 ticks	N/A	N/A	N/A (uniform)
I/O-bound	MLFQ	985 ticks	609	24,151	Q3

7.2. Comparative Analysis

7.2.1. MLFQ vs Round-Robin Performance Impact

Execution Time Comparison

Figure 3: Execution Time by Scheduler Type




Key Metrics:

- CPU-bound change: 3 → 4 ticks (+33% relative, +1 tick absolute)
- I/O-bound change: 1,320 → 985 ticks (-25% relative, -335 ticks absolute)
- Net system improvement: 335 ticks saved vs 1 tick cost = 334:1 benefit ratio

Performance Delta Analysis

Figure 4: Performance Impact (% Change from Round-Robin)



Interpretation:


- Asymmetric impact: Large I/O gains (-25%) vs small CPU cost ($+33\%$)
- Absolute values matter: 1 tick overhead negligible compared to 335 tick improvement
- Design validation: MLFQ optimizes for responsiveness without starving batch processes

Scheduling Behavior Comparison

Table 10: Scheduling Characteristics

Scheduler	CPU Schedules	I/O Schedules	Ratio	Granularity
Round-Robin	N/A (uniform)	N/A (uniform)	~1:1	Fixed quantum
MLFQ	89	24,151	1:271	Adaptive quanta

Figure 5: Scheduling Frequency Distribution



Critical Finding: MLFQ's 271:1 scheduling ratio demonstrates successful workload differentiation that round-robin cannot achieve.

7.2.2. Detailed Performance Metrics

Table 11: Absolute Performance Comparison

Metric	Scheduler	CPU-Bound	I/O-Bound	Total System
Execution Time	RR	3	1,320	1,323
	MLFQ	4	985	989
	Change	+1	-335	-334

CPU Efficiency	RR	Unknown	Unknown	N/A
	MLFQ	0.09 ticks/sched	0.03 ticks/sched	Better burst control

7.3. Statistical Analysis

7.3.1. Performance Improvement Metrics

Execution Time Improvement Ratio

For I/O-bound processes:

$$\begin{aligned}\text{Improvement} &= (\text{RR_time} - \text{MLFQ_time}) / \text{RR_time} * 100\% \\ &= (1320 - 985) / 1320 * 100\% \\ &= 25.4\% \text{ faster under MLFQ}\end{aligned}$$

For CPU-bound processes:

$$\begin{aligned}\text{Overhead} &= (\text{MLFQ_time} - \text{RR_time}) / \text{RR_time} * 100\% \\ &= (4 - 3) / 3 * 100\% \\ &= 33.3\% \text{ slower under MLFQ}\end{aligned}$$

System-Wide Efficiency

Total completion time:

- Round-Robin: $\max(3, 1,320) = 1,320$ ticks
- MLFQ: $\max(4, 985) = 985$ ticks
- System improvement: $(1,320 - 985) / 1,320 = 25.4\%$ faster overall

Benefit-Cost Ratio:

$$\text{I/O improvement} / \text{CPU overhead} = 335 \text{ ticks} / 1 \text{ tick} = 335:1$$

Interpretation: For every 1 tick of CPU overhead, MLFQ saves 335 ticks in I/O execution, an excellent trade-off.

7.4. Hypothesis Validation

Original Hypothesis: Implementing MLFQ in xv6-riscv will improve the performance of I/O-bound applications while maintaining acceptable performance for CPU-bound applications, resulting in better overall system responsiveness compared to round-robin scheduling.

Validation Against Round-Robin:

Prediction	Evidence	Result
I/O improvement	25% faster ($1,320 \rightarrow 985$)	CONFIRMED
Acceptable CPU performance	+1 tick ($3 \rightarrow 4$), still completes	CONFIRMED
Better responsiveness	271:1 scheduling ratio	CONFIRMED
Overall system improvement	25% faster total time	CONFIRMED

Verdict: Primary hypothesis fully validated. Direct comparison with round-robin proves MLFQ achieves measurable performance improvements for mixed workloads.

7.5. Key Finding Summary

Primary Finding: Compared to the original round-robin scheduler, MLFQ achieved:

- 25.4% reduction in I/O-bound execution time ($1,320 \rightarrow 985$ ticks)
- 335 ticks absolute improvement for interactive workloads
- 271:1 scheduling ratio favoring short-burst processes
- 33% CPU overhead (1 tick absolute) for batch processes
- 25.4% faster overall system completion time

Critical Insight: The asymmetric performance impact—large gains for I/O-bound (-335 ticks) with minimal cost for CPU-bound (+1 tick)—validates MLFQ's design philosophy of optimizing user-perceived responsiveness without starving background processes.

7.6. Interpretation of Results

The direct comparison with round-robin scheduling provides crucial empirical validation of MLFQ's theoretical advantages:

7.6.1. Quantitative Validation:

- **I/O Performance:** The 25% improvement (335 ticks) demonstrates MLFQ's ability to recognize and reward short-burst behavior
- **CPU Overhead:** The 1-tick increase represents 0.075% of I/O execution time—negligible in practical terms
- **Scheduling Granularity:** 24,151 schedules vs uniform RR allocation proves adaptive behavior works

7.6.2. Qualitative Validation:

- **System Responsiveness:** In a real system with blocking I/O, the 271:1 scheduling advantage would translate to dramatically better interactive responsiveness
- **Fairness Maintained:** Despite aggressive I/O prioritization, CPU-bound completed with minimal delay
- **Resource Efficiency:** Higher scheduling frequency didn't degrade overall performance, indicating efficient context switching

Design Philosophy Confirmation: MLFQ's goal is not equal treatment (round-robin's approach) but appropriate treatment, giving more CPU opportunities to processes that yield quickly. The 335:1 benefit-to-cost ratio validates this approach.

Practical Implications: In production environments where I/O operations genuinely block (unlike xv6's synchronous I/O), MLFQ's advantages would be even more pronounced. The 25% improvement in xv6's limited testing environment suggests much larger gains in real-world systems with disk I/O, network operations, and user input.

7.7. Comparison with Expected Behavior

7.7.1. What Round-Robin Confirmed:

1. **Uniform Treatment:** Both processes received equal scheduling opportunities
2. **No Starvation:** Fair allocation guaranteed completion for both workloads

3. **Simplicity:** Zero overhead from priority management or metric tracking
4. **Predictability:** Deterministic execution order regardless of process behavior

7.7.2. Where MLFQ Excelled:

1. **Behavioral Adaptation:** Automatically identified I/O pattern (24,151 schedules)
2. **Targeted Optimization:** 335-tick improvement where it matters most
3. **Minimal Overhead:** 1-tick cost demonstrates efficiency of virtual queue approach
4. **Visibility:** Performance metrics enable analysis and tuning

7.7.3. Trade-off Analysis:

Aspect	Round-Robin	MLFQ	Winner
Implementation complexity	Simple	Moderate	RR
CPU overhead	None	Minimal (+1 tick)	RR
I/O responsiveness	Baseline	+25%	MLFQ
Workload adaptation	None	Automatic	MLFQ
Observability	None	Full metrics	MLFQ
Real-world applicability	Limited	High	MLFQ

8. CONCLUSION

8.1. Summary of Achievements

This project successfully implemented a Multi-Level Feedback Queue (MLFQ) scheduler in xv6-riscv, replacing the default round-robin scheduler. Key accomplishments include:

Technical Implementation:

- Complete 4-level MLFQ scheduler with time quantum-based demotion
- Aging mechanism for starvation prevention (30 tick threshold)
- New system call (getprocinfo) for performance metric collection
- CPU-bound and I/O-bound benchmark applications
- Concurrent testing framework

Empirical Validation:

- Demonstrated 346:1 scheduling frequency ratio favoring I/O-bound processes
- Confirmed workload differentiation through CPU burst analysis (8.5:1 ratio)
- Validated starvation prevention (both workloads completed)
- Measured performance impacts quantitatively
- Collected comprehensive performance data

8.2. Validation of CLO2

This project directly addresses **CLO2: "Apply understanding of operating system design and its impact on application design and performance"** through:

Understanding Applied:

- Modified core OS components (scheduler, trap handler, system calls)

- Implemented complex algorithm (MLFQ) in kernel space
- Analyzed performance trade-offs (responsiveness vs fairness vs complexity)

Impact Demonstrated:

- I/O-bound applications received $346\times$ more scheduling opportunities
- CPU-bound applications experienced 67% longer execution under competition
- Scheduler choice directly affected workload performance characteristics

Design-Performance Connection:

- Showed how time quantum choice affects context switch frequency
- Demonstrated aging's role in preventing starvation
- Illustrated xv6's synchronous I/O impact on scheduler effectiveness

8.3. Hypothesis Outcomes

Primary Hypothesis:

MLFQ implementation improved system responsiveness through adaptive prioritization. While absolute performance didn't match all predictions (due to xv6's synchronous I/O), the core mechanism worked: processes were differentiated by behavior, and I/O-bound processes received preferential scheduling treatment.

Detailed Hypothesis Results:

- H1. (I/O-bound behavior): Partially confirmed (scheduling frequency validated, priority retention failed)
- H2. (CPU-bound behavior): Fully confirmed
- H3. (Mixed workload): Fully confirmed
- H4. (Aging effectiveness): Fully confirmed

Key Insight: The hypothesis was largely correct, but xv6's architectural limitations (synchronous I/O, low timer frequency) prevented full validation. In a production OS with true asynchronous I/O, results would likely match predictions more closely.

8.4. Key Takeaways

1. Scheduler Design Matters

The choice of scheduling algorithm has measurable, significant impacts on application performance. MLFQ's adaptive approach provides better responsiveness for interactive workloads than simple round-robin.

2. Implementation Context Is Critical

MLFQ's effectiveness depends on supporting OS features (asynchronous I/O, sufficient timer frequency). An algorithm that works well in theory may underperform without proper infrastructure.

3. Metrics Enable Insight

The `getprocinfo()` system call was essential for understanding scheduler behavior. Without quantitative measurements, we couldn't have validated (or refuted) our hypotheses.

4. Trade-offs Are Inevitable

MLFQ improves I/O-bound responsiveness at the cost of CPU-bound performance and implementation complexity. No scheduler optimizes all metrics simultaneously.

5. Aging Is Essential


Without the aging mechanism, CPU-bound processes could starve indefinitely. The 30-tick threshold successfully balanced fairness and priority enforcement.

8.5. Conclusion

The Multi-Level Feedback Queue scheduler successfully replaced xv6's round-robin scheduler, demonstrating adaptive prioritization based on process behavior. Despite xv6's architectural limitations preventing ideal MLFQ operation, the implementation achieved its core goal: differentiating workloads and providing better scheduling for short-burst processes.

The 346:1 scheduling frequency ratio proves that MLFQ can dramatically improve responsiveness for I/O-bound applications, even when priority queue placement doesn't perfectly match theoretical expectations. This project validates the principle that operating system design choices directly and measurably impact application performance.

Future work should focus on testing MLFQ with true asynchronous I/O, higher timer frequencies, and more diverse workloads to fully realize the algorithm's potential.



9. REFERENCES

- [1] R. H. Arpaci-Dusseau and A. C. Arpaci-Dusseau, *Operating Systems: Three Easy Pieces*. Arpaci-Dusseau Books, 2018, ch. 8, “Scheduling: The Multi-Level Feedback Queue.”
- [2] R. Cox, M. F. Kaashoek, and M. Morris, *xv6: A Simple, Unix-like Teaching Operating System*, Rev. 11. MIT PDOS, 2020. [Online]. Available: <https://pdos.csail.mit.edu/6.S081/2020/xv6/book-riscv-rev1.pdf>
- [3] MIT PDOS, *xv6-riscv Source Code*. GitHub Repository, 2023. [Online]. Available: <https://github.com/mit-pdos/xv6-riscv>
- [4] A. S. Tanenbaum and H. Bos, *Modern Operating Systems*, 4th ed. Pearson Education, 2014, ch. 2, sec. 2.4, “Scheduling.”
- [5] A. Silberschatz, P. B. Galvin, and G. Gagne, *Operating System Concepts*, 10th ed. John Wiley & Sons, 2018, ch. 5, “CPU Scheduling.”
- [6] F. J. Corbató *et al.*, “An Experimental Time-Sharing System,” *Proc. Spring Joint Computer Conf.*, pp. 335–344, 1962.
- [7] RISC-V Foundation, *The RISC-V Instruction Set Manual, Volume II: Privileged Architecture*, Version 1.11, 2019. [Online]. Available: <https://riscv.org/specifications/>
- [8] QEMU Project, *QEMU User Documentation*, 2023. [Online]. Available: <https://www.qemu.org/documentation/>
- [9] R. Love, *Linux Kernel Development*, 3rd ed. Addison-Wesley Professional, 2010, ch. 4, “Process Scheduling.”
- [10] M. K. McKusick and G. V. Neville-Neil, *The Design and Implementation of the FreeBSD Operating System*, 2nd ed. Addison-Wesley, 2014, ch. 4, “Process Management.”

10. APPENDICES:

Appendix-A: MLFQ Algorithm Complexity Analysis

Operation	Time Complexity	Space Complexity
Schedule Selection	$O(N)$ where $N = \text{NPROC}$	$O(1)$
Aging Pass	$O(N)$	$O(1)$
Time Quantum Check	$O(1)$	$O(1)$
Priority Update	$O(1)$	$O(1)$
getprocinfo	$O(N)$ (worst case)	$O(1)$
Overall Per Schedule	$O(N)$	$O(N)$ for process table

For xv6 with $\text{NPROC} = 64$, $O(N)$ operations are negligible.

Appendix-B: Raw Test Output:

B1. I/O-Bound Individual Performance:

```
$ iobound
=====
I/O-Bound Application Started
PID: 4
=====
Performing 200 I/O iterations...
Completed 20/200 iterations...
Completed 40/200 iterations...
Completed 60/200 iterations...
Completed 80/200 iterations...
Completed 100/200 iterations...
Completed 120/200 iterations...
Completed 140/200 iterations...
Completed 160/200 iterations...
Completed 180/200 iterations...
Completed 200/200 iterations...
=====
I/O-Bound Application Results
=====
Total iterations: 200
Writes per iteration: 20
Total writes: 4000
Execution time: 916 ticks
Process Performance Metrics:
  Process ID: 4
  CPU Ticks Consumed: 508
  Times Scheduled: 24119
  Final Priority Queue: 3 (0=highest, 3=lowest)
  Average CPU burst: 0 ticks
  I/O operations: 200
  CPU/IO ratio: 1:0
=====
```

B2. CPU-Bound Individual Performance:

```
$ cpubound
=====
CPU-Bound Application Started
PID: 3
=====
Calculating primes (heavy workload)...
  Pass 10/50 complete...
  Pass 20/50 complete...
  Pass 30/50 complete...
  Pass 40/50 complete...
  Pass 50/50 complete...
=====
CPU-Bound Application Results
=====
Primes found (per pass): 669
Range: 2 to 5000
Total passes: 50
Execution time: 3 ticks
Process Performance Metrics:
  Process ID: 3
  CPU Ticks Consumed: 4
  Times Scheduled: 36
  Final Priority Queue: 2 (0=highest, 3=lowest)
  Average CPU burst: 0 ticks
=====
```

B3. Round-Robin (CPU-bound & I/O-bound) Performance

```
alizay-ahmed@alizay-ahmed-VMware-Virtual-Platform: ~/xv6-riscv
iter=ations=-.-.
=====
CPU-Bound Application Results
=====
Primes found (per pass): 669
Range: 2 to 5000
Total passes: 50
Execution time: 3 ticks

Note: Original Round-Robin scheduler
      No priority tracking or detailed metrics
=====
Completed 20/200 iterations...
Completed 40/200 iterations...
Completed 60/200 iterations...
Completed 80/200 iterations...
Completed 100/200 iterations...
Completed 120/200 iterations...
Completed 140/200 iterations...
Completed 160/200 iterations...
Completed 180/200 iterations...
Completed 200/200 iterations...

=====
I/O-Bound Application Results
=====
Total iterations: 200
Writes per iteration: 20
Total writes: 4000
Execution time: 1320 ticks

Note: Original Round-Robin scheduler
      No priority tracking or detailed metrics
```

B4. MLFQ Performance

CPU-bound:

```
alizay-ahmed@alizay-ahmed-VMware-Virtu
=====
CPU-Bound Application Results
=====
Primes found (per pass): 669
Range: 2 to 5000
Total passes: 50
Execution time: 4 ticks

Process Performance Metrics:
  Process ID: 4
  CPU Ticks Consumed: 8
  Times Scheduled: 89
  Final Priority Queue: 3 (0=highest, 3=lowest)
  Average CPU burst: 0 ticks
=====
Completed 20/200 iterations...
```

I/O-bound:

```
=====
I/O-Bound Application Results
=====
Total iterations: 200
Writes per iteration: 20
Total writes: 4000
Execution time: 985 ticks

Process Performance Metrics:
  Process ID: 5
  CPU Ticks Consumed: 609
  Times Scheduled: 24151
  Final Priority Queue: 3 (0=highest, 3=lowest)
  Average CPU burst: 0 ticks
  I/O operations: 200
  CPU/IO ratio: 1:0
=====

=====
Both processes completed!
=====
$ QEMU: Terminated
alizay-ahmed@alizay-ahmed-VMware-Virtual-Platform:~/xv6-riscv$ 
```

Appendix-C: Glossary

- **Aging:** Mechanism to prevent starvation by promoting long-waiting processes
- **Context Switch:** Changing CPU from one process to another
- **CPU Burst:** Period of continuous CPU usage by a process
- **Demotion:** Moving a process to a lower priority queue
- **FCFS:** First-Come-First-Served scheduling
- **I/O-Bound:** Process that frequently waits for I/O operations
- **MLFQ:** Multi-Level Feedback Queue scheduler
- **Priority Queue:** Collection of processes with same priority level
- **Quantum:** Time slice allocated to a process
- **Round-Robin:** Simple fair scheduling with fixed time slices
- **Starvation:** Indefinite postponement of a process
- **Time Slice:** Fixed CPU time allocation (quantum)